

Exercise 5

Power Method, BLAS/LAPACK, OpenMP

Solutions review & feedback


Question 1a: Power method

```
#include <random>

// alpha: Diagonal scaling factor
// A:      N x N Matrix
// N:      Matrix dimension
void initialize_matrix(const double alpha,
                      double* const A, const int N)
{
    // random generator with seed 0
    std::default_random_engine g(0);
    // uniform distribution in [0, 1]
    std::uniform_real_distribution<double> u;

    // matrix initialization
    for (int i = 0; i < N; i++) {
        for (int j = i+1; j < N; j++) {
            const double rand = u(g);
            A[i*N + j] = rand;
            A[j*N + i] = rand;
        }
        A[i*N + i] = (i+1.0) * alpha;
    }
}
```

1. Symmetric matrix, A:
 $A[i,j] = A[j,i]$



Question 1a: Power method

...

```
auto tstart = std::chrono::steady_clock::now();
// Main algorithm:
// You can re-structure the Power Method algorithm such that you only
// perform one GEMV operation in the iteration loop. By doing so, you have
// to use a pointer swap at the end of the loop (which is very cheap to
// do).
_gemv(N, N, A, q0, q1);
while (true)
{
    _norm(N, q1);
    _gemv(N, N, A, q1, q0);
    lambda1 = 0.0;
    for (int i = 0; i < N; ++i)
        // estimate eigenvalue with Rayleigh quotient
        lambda1 += q0[i] * q1[i];
    ++iter;
    if (std::abs(lambda1 - lambda0) < tol)
        break;
    lambda0 = lambda1;
    std::swap(q0, q1);
}
// end of algorithm
auto tend = std::chrono::steady_clock::now();
auto time = std::chrono::duration_cast<std::chrono::milliseconds>(tend -
    tstart).count();
```

...

1. Timing should include first gemv since it is part of the algorithm
2. Single gemv in iteration loop
3. No need to repeat gemv calculation to compute eigenvalue!

$$\mathbf{q}^{(k+1)} = \frac{A\mathbf{q}^{(k)}}{\|A\mathbf{q}^{(k)}\|_2}$$

$$\lambda = \frac{\mathbf{q}^T A \mathbf{q}}{\mathbf{q}^T \mathbf{q}}$$

4. Use pointer swap to update eigenvectors (or equivalent design)

Question 1b+c: Power method (CBLAS, LAPACK)

CBLAS:

```
...
auto tstart = std::chrono::steady_clock::now();
// Main algorithm:
// You can re-structure the Power Method algorithm such that you only
// perform one GEMV operation in the iteration loop. By doing so, you have
// to use a pointer swap at the end of the loop (which is very cheap to
// do).
cbblas_dsymv(CblasRowMajor, CblasUpper, N, 1.0, A, N, q0, 1, 0.0, q1, 1);
while (true)
{
    const double norm = cbblas_dnorm2(N, q1, 1);
    cbblas_dscal(N, 1./norm, q1, 1);
    cbblas_dsymv(CblasRowMajor, CblasUpper, N, 1.0, A, N, q1, 1, 0.0, q0, 1);
    lambda1 = cbblas_ddot(N, q0, 1, q1, 1);
    ++iter;
    if (std::abs(lambda1 - lambda0) < tol)
        break;
    lambda0 = lambda1;
    std::swap(q0, q1);
}
// end of algorithm
auto tend = std::chrono::steady_clock::now();
auto time = std::chrono::duration_cast<std::chrono::milliseconds>(tend -
    tstart).count();
...
```

1. Same matrix-vector and vector operations, but replaced by CBLAS function calls
2. Lecture slides: usage example for CBLAS function calls
3. Use of `std::swap(...)` instead of `cbblas_dcopy(...)`

LAPACK:


```
...
auto tstart = std::chrono::steady_clock::now();
int info = LAPACKE_dsyev(LAPACK_ROW_MAJOR, jobz, uplo, N, A, N, lambdas);
auto tend = std::chrono::steady_clock::now();
auto time = std::chrono::duration_cast<std::chrono::milliseconds>(tend -
    tstart).count();
...
```

Single function call

Question 1b+c: Power method (CBLAS, LAPACK)



CBLAS

difference in time-to-solution spans 2 orders of magnitude



LAPACK

time-to-solution invariant of α




- Power method: iterative algorithm
- Might not always be the best choice - this depends on the properties of the matrix we are interested in!


Question 1d: Convergence of the power method



Number of iterations



Eigenvalue ratio



Question 2: OpenMP bug hunts

```
...
10 #pragma omp parallel
11 {
12
13     for (int step=0; step<100; step++)
14     {
15         #pragma omp parallel for nowait
16         for (int i=1; i<n; i++) {
17             b[i-1] = (a[i]+a[i-1])/2.;
18             c[i-1] += a[i];
19         }
20
21         #pragma omp for
22         for (int i=0; i<m; i++)
23             z[i] = sqrt(b[i]+c[i]);
24
25         #pragma omp for reduction(+:sum)
26         for (int i=0; i<m; i++)
27             sum = sum + z[i];
28
29         #pragma omp critical
30         {
31             do_work(t, sum);
32         }
33
34         #pragma omp single
35         {
36             t = new_value(step);
37         }
38     }
39 }
...
```

Remove both!



Same iteration space! Combination in single for loop to save the overhead of one #pragma omp for

critical is unnecessary. Remove and add a barrier after function do_work.

Explicitly state the assumption that do_work can be executed by a single thread since it's arguments are constant and the function computes a total sum. In this case critical should be replaced with single.

See solution pdf for full solution of Question 2

General remarks

- 👍 Avoid copy-pasting
- 👍 Do submit your codes along with pdf solutions and results
- 👍 Collect results (text, figures, theoretical proofs in a pdf file)
- 👍 If in doubt, google or ask on piazza!

e.g.



power_manual.cpp



power_cblas.cpp



eigenv_lapack.cpp

```
#include "common.h"
```

```
...
```

```
// matrix initialization  
double* A = new double[N*N];  
initialize_matrix(alpha, A, N);
```



common.h

```
void initialize_matrix(const double alpha,  
                      double* const A,  
                      const int N);
```