


Exercise 7

High Performance Computing for Science and Engineering

Guido Novati

November 10, 2017


Convolutional Neural Networks



LeNet of Yann LeCun et al., 1998

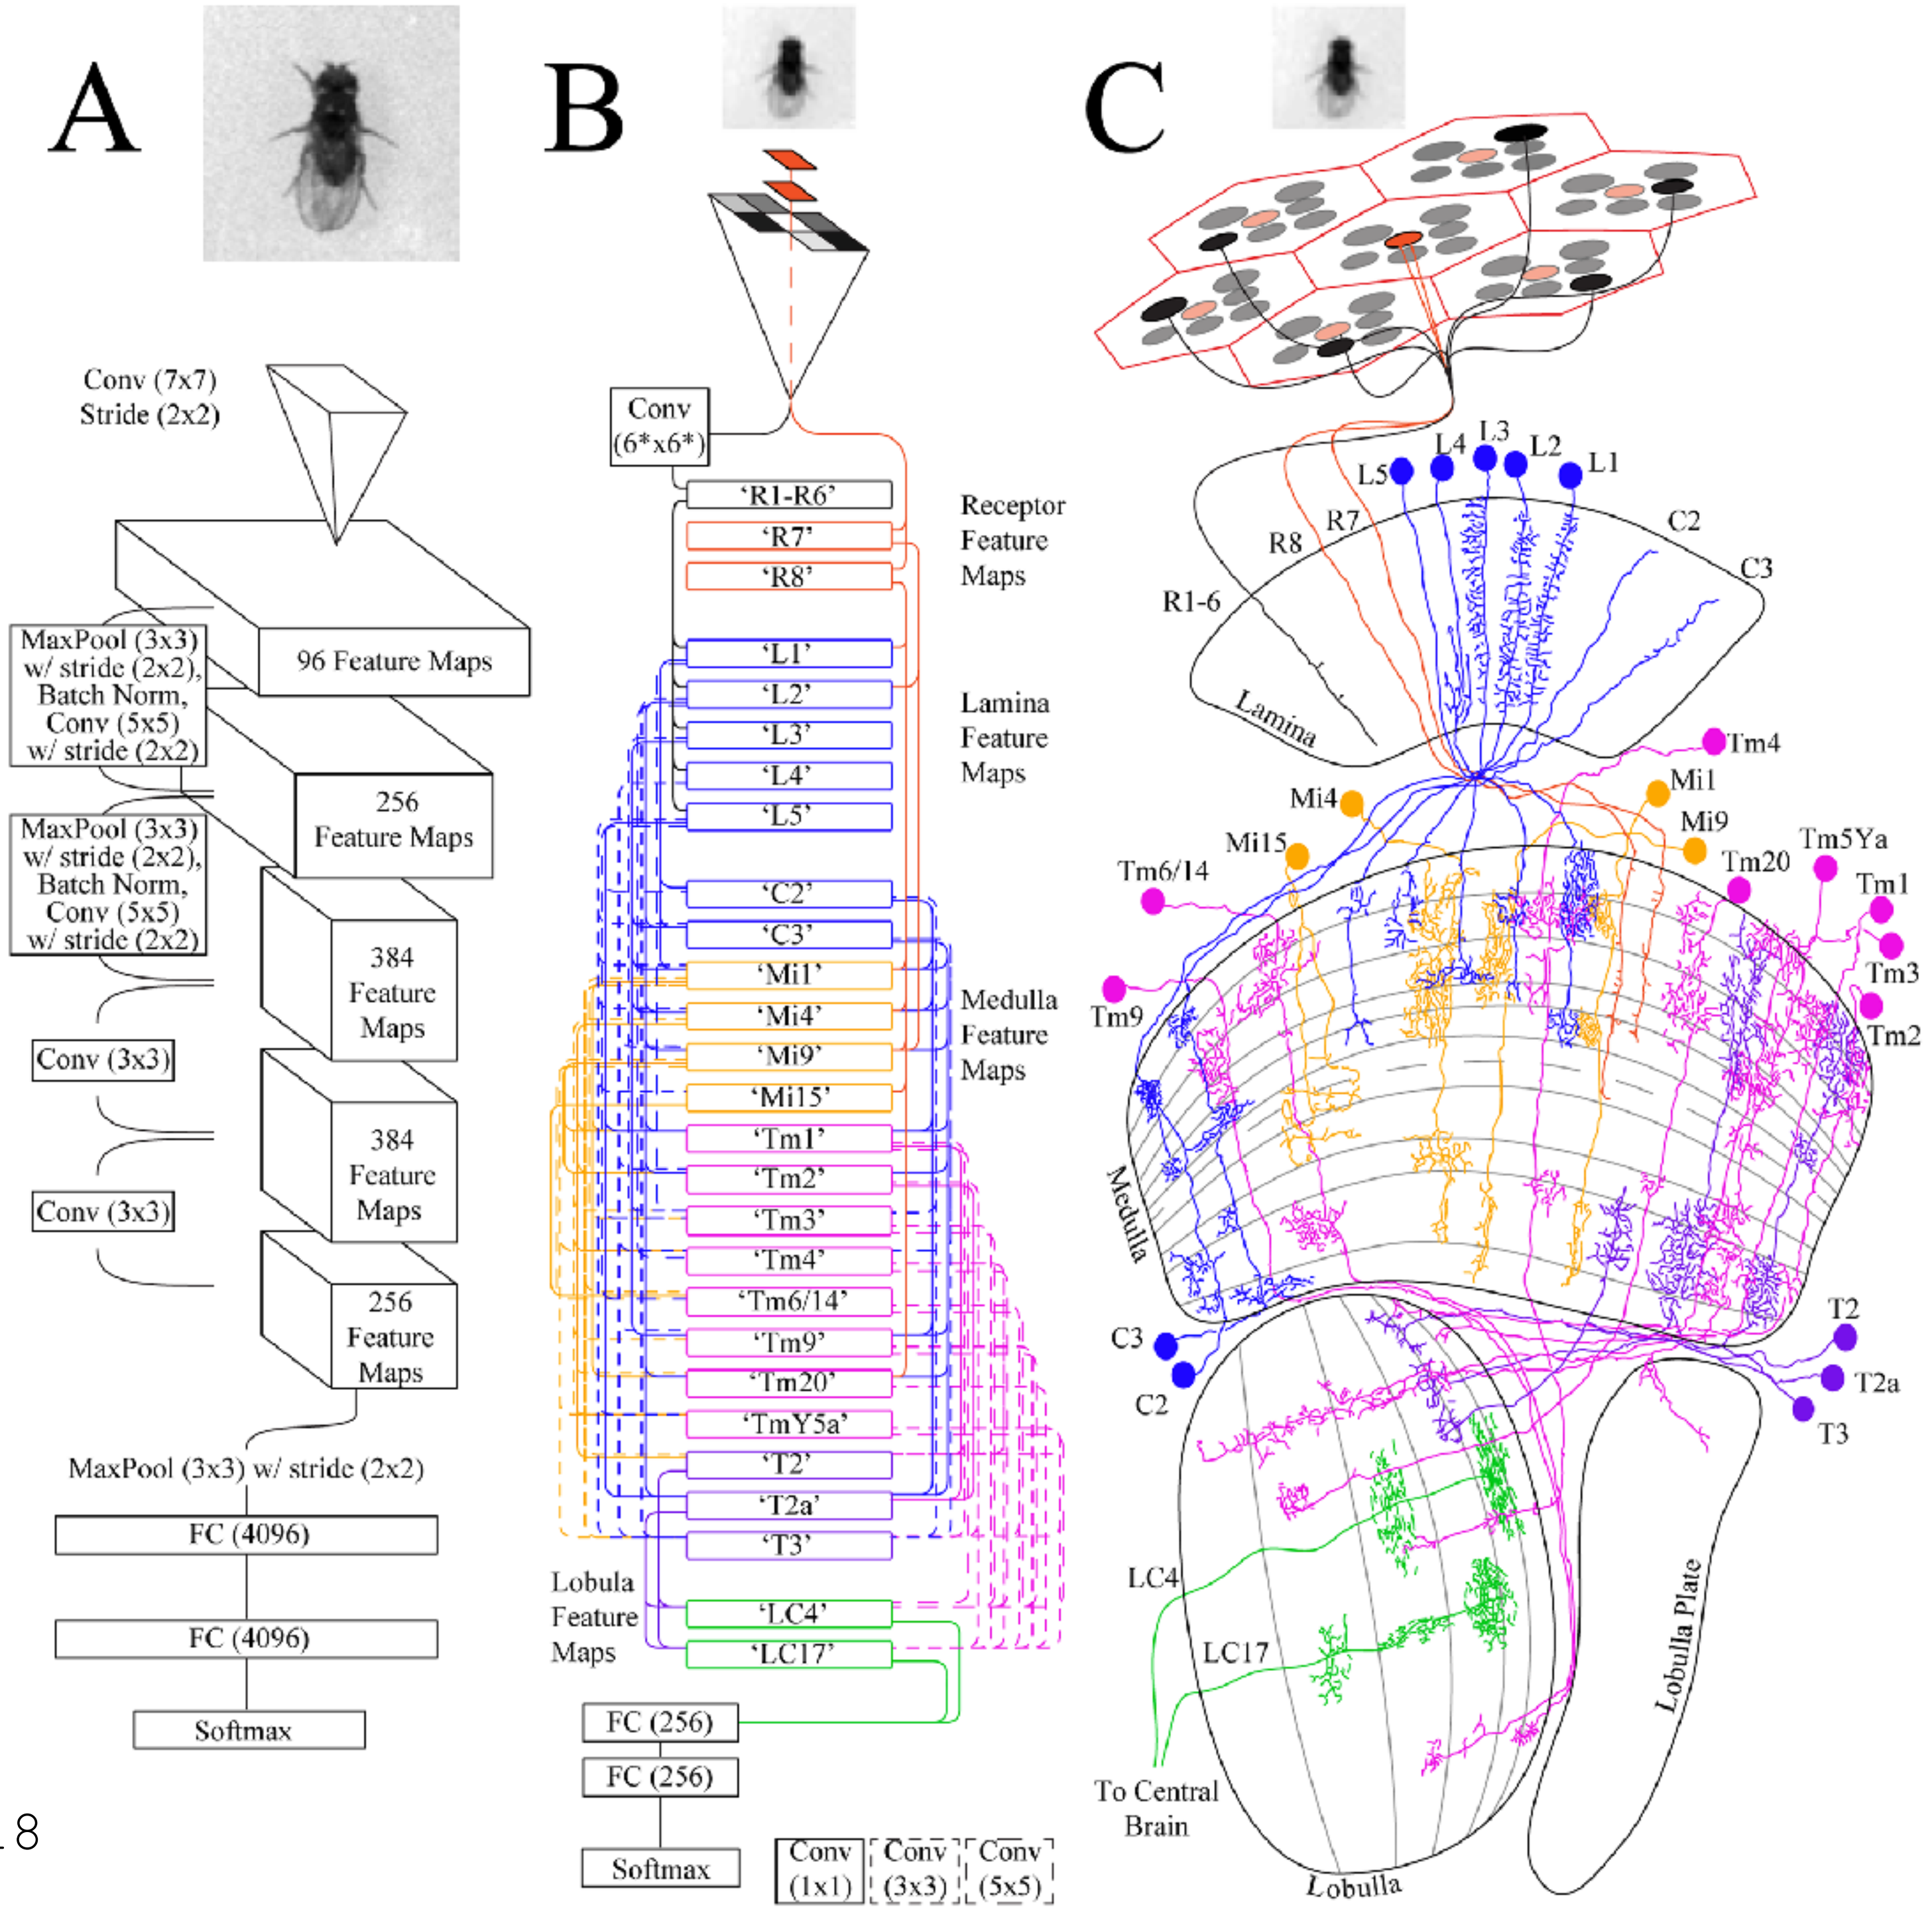
- Before deep learning, AI meant expertly designed “if” statements
- With deep learning, AI means GEMM.
- We (HPC people) are pretty good at doing GEMM.
- CNN are the backbone of recent hype in “deep learning”
- Parametric models that are well suited to classify / recognise image contents

Biological Intuition



- **Very** roughly speaking, biological brains have neurons that activate when they recognize a triggering pattern in their inputs.
- Each unit does “simple” pattern recognition. Complexity emerges from sheer numbers

Convolutional model of **a part of** the Fruit-fly's brain




Jonathan Schneider et al., 2018

What is a Convolution (in machine learning) (1)

- A convolution is a parametric operation that maps an “image” to an other “image”


Input is a matrix :

$$\text{InY} * \text{InX} * \text{InC}$$




Parameters are a matrix:

$$\text{KnY} * \text{KnX} * \text{InC} * \text{KnC}$$



Output is a matrix :


$$\text{OpY} * \text{OpX} * \text{KnC}$$



- The third dimension are the “color” channels (or feature maps)
- InC and KnC can be any number
- Parameters are called “filters” or “kernels”


What is a Convolution (in machine learning) (2)

- Convolution is performed by iterating along X and Y
- At each position, and for each filter, we compute the scalar product between the filter and a patch of the image ($K_n Y * K_n X * I_n C$)
- The output of the scalar prod is a number which is written onto one color pixel of the output image




1D, 1 Filter

- Let's apply convolution in 1D:




1D, 1 Filter

- Let's apply convolution in 1D:



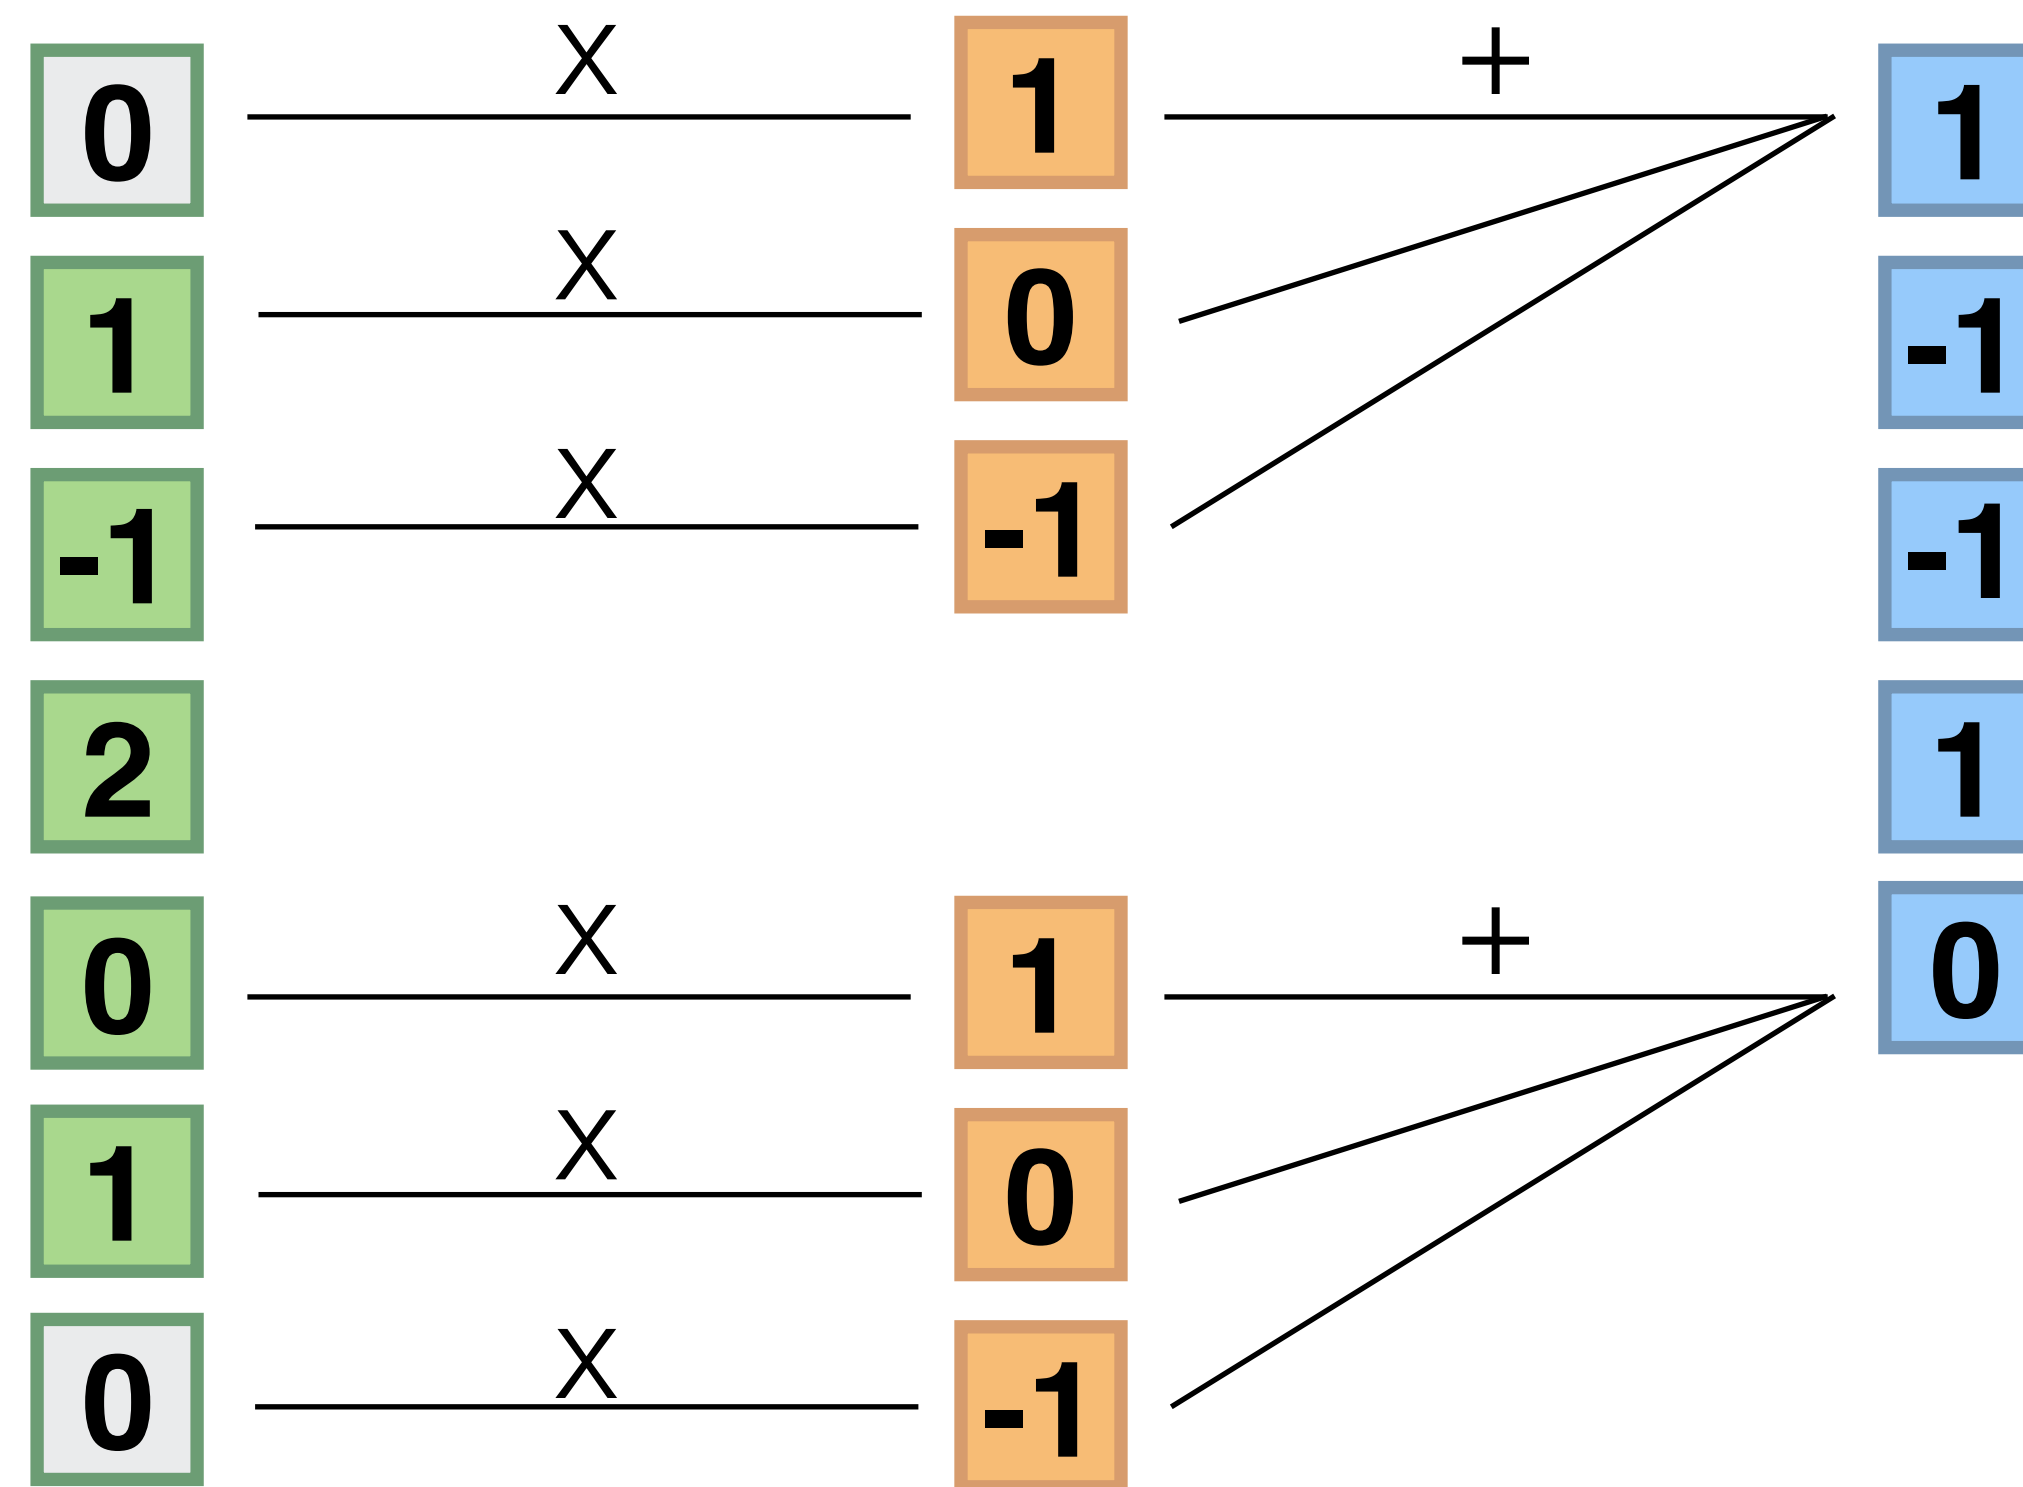
1D, 1 Filter

- Let's apply convolution in 1D:




1D, 1 Filter

- Padding:
 - Size of the image is extended in both directions by P_y and P_x
 - Padding pixels usually carry value 0 (“zero padding”)



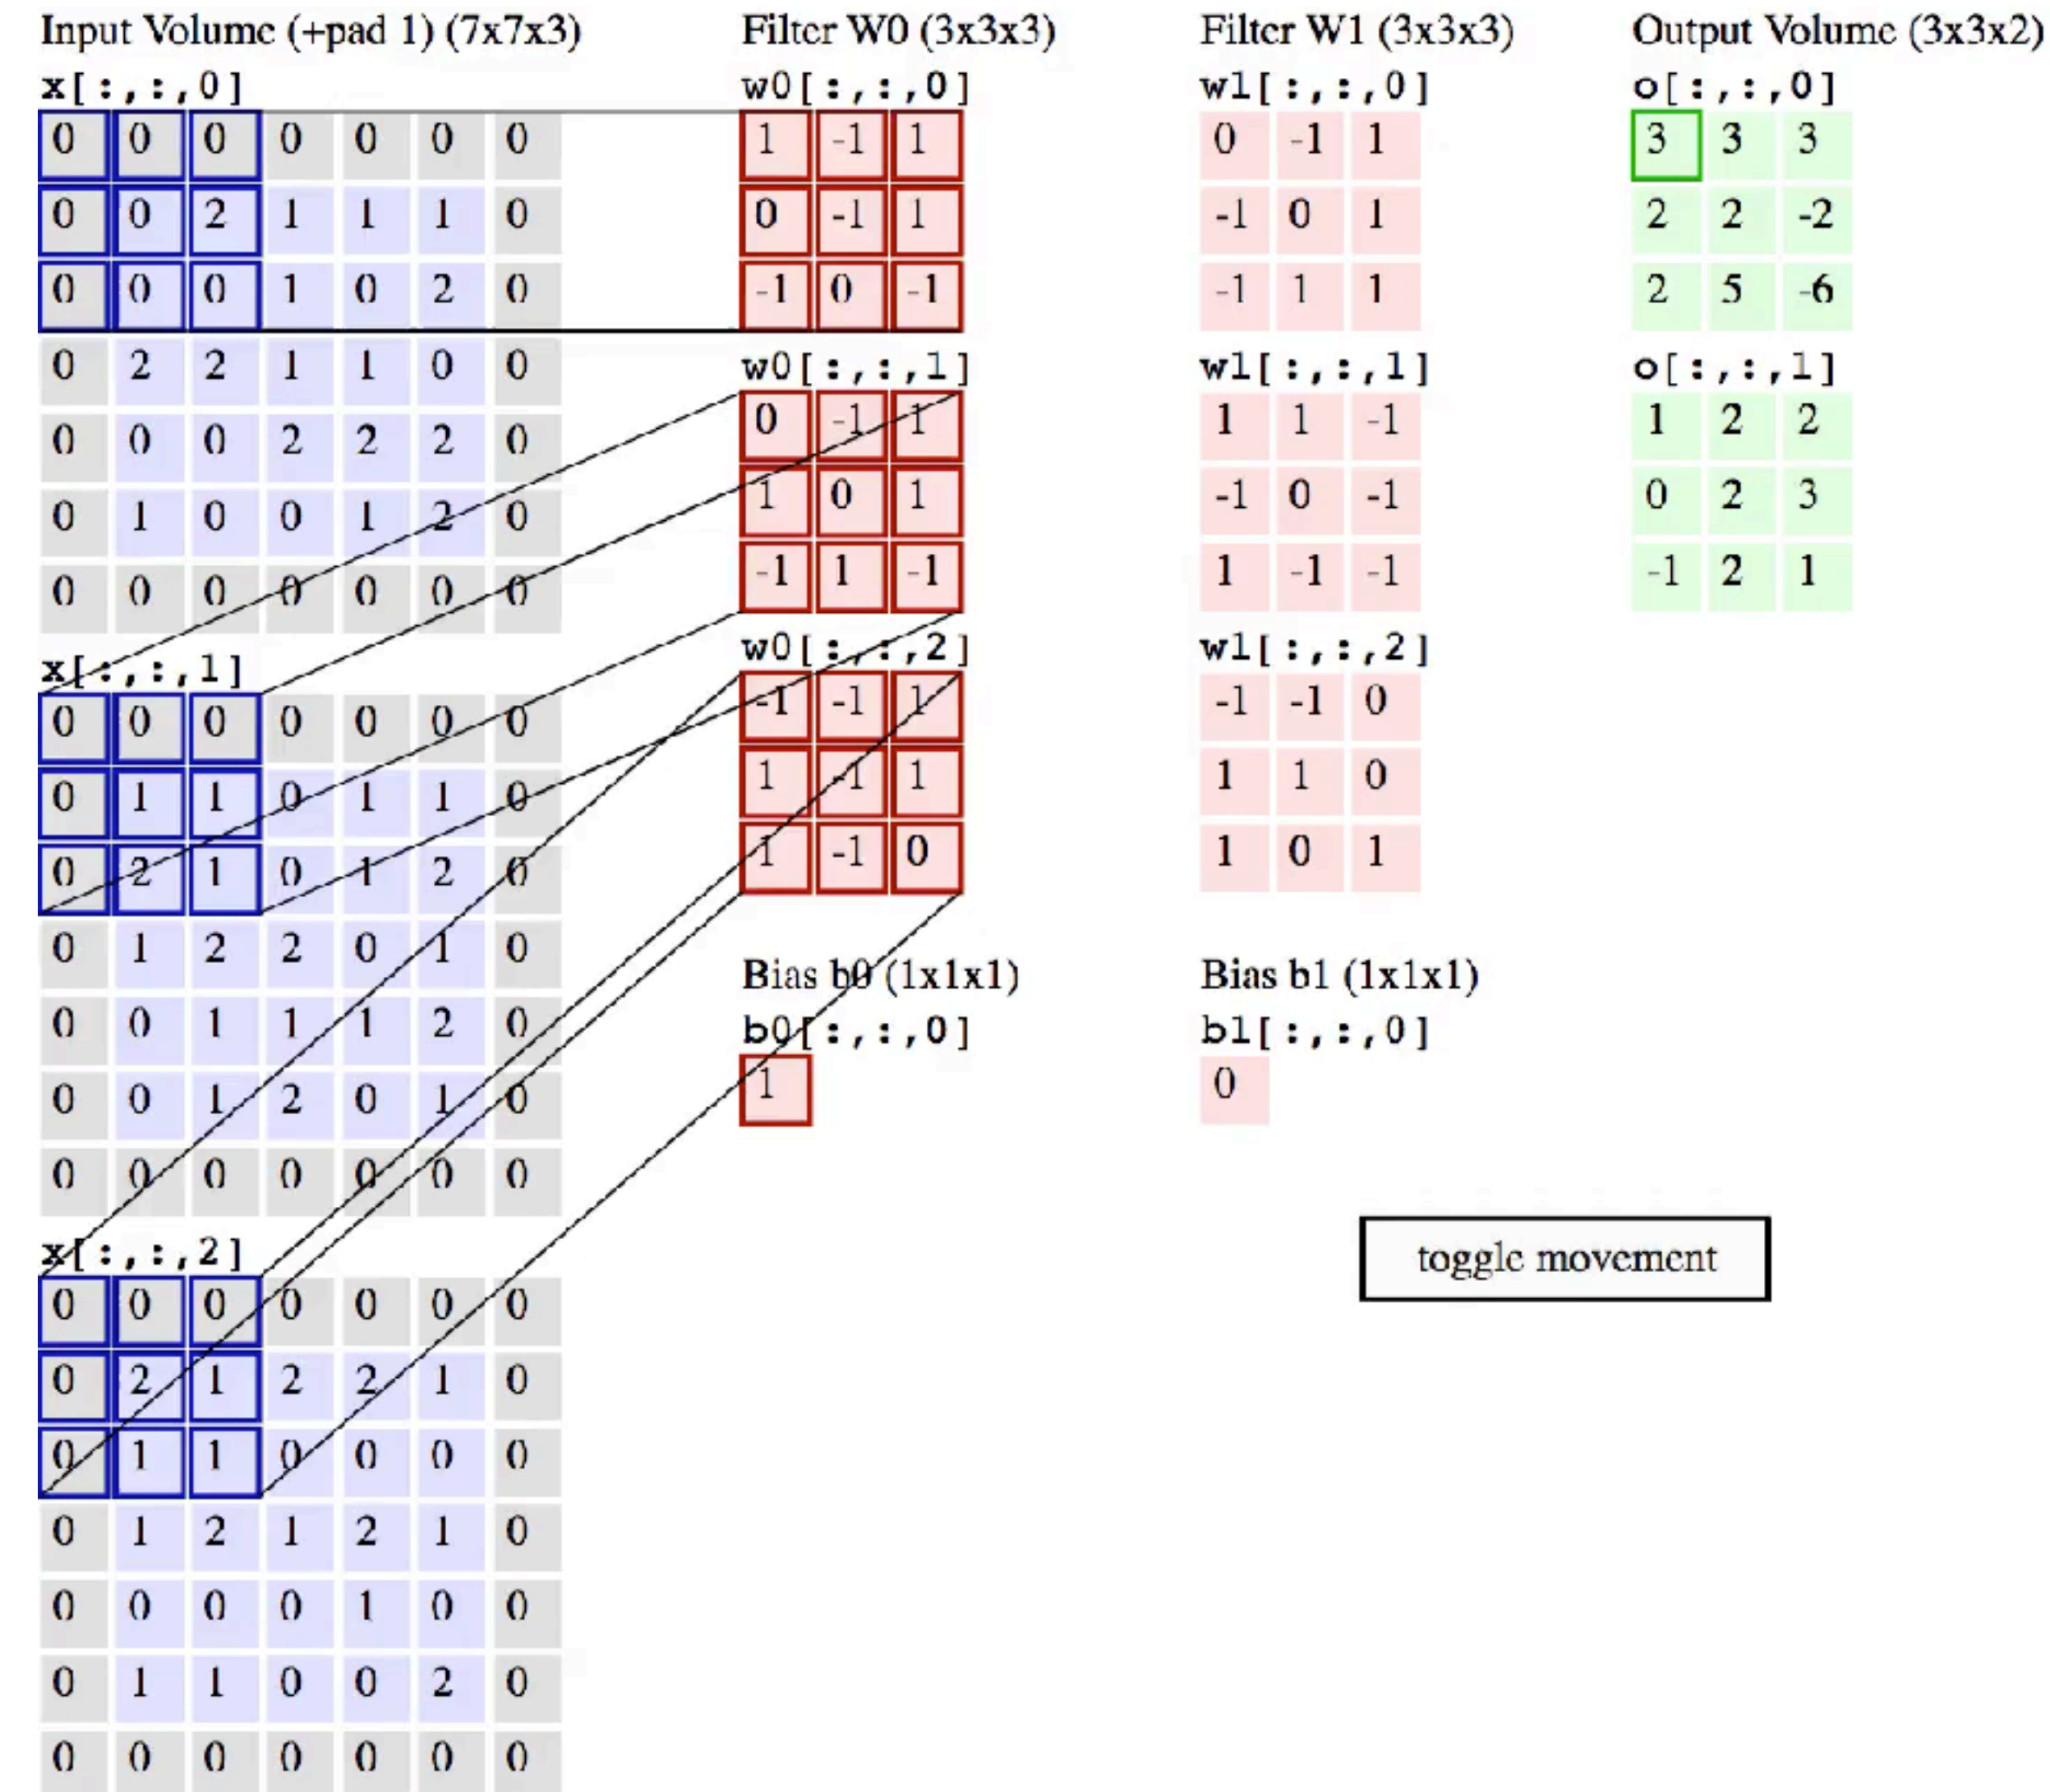
1D, 1 Filter

- Stride:
- Convolution does not need to be computed by iterating with increments of 1 pixel
- Here: padding 1, stride 2:




2D convolutions rules of thumb

- If the input image is $\text{InY} * \text{InX} * \text{InC}$
- The Filters are $\text{KnY} * \text{KnX} * \text{InC} * \text{KnC}$
- With strides S_y and S_x
- And padding P_y and P_x
- The output image has size:
 - $\text{OpY} = (\text{InY} - \text{KnY} + 2P_y) / S_y + 1$
 - $\text{OpX} = (\text{InX} - \text{KnX} + 2P_x) / S_x + 1$
 - KnC
- These numbers must be integers!
This requirement may force you to tweak the parameters of the convolution



2D convolution forward loop:

- forward operation requires iterating through input and computing the inner product




```
for (int oy = 0; oy < OpY; oy++) // loop over output image Y
for (int ox = 0; ox < OpX; ox++) // loop over output image X
{
    for (int fy = 0; fy < KnY; fy++) // loop over filter Y
    for (int fx = 0; fx < KnX; fx++) // loop over filter X
    {
        //index along input map of the convolution op:
        const int ix = ox * Sx - Px + fx;
        const int iy = oy * Sy - Py + fy;
        //padding: skip addition if outside input boundaries
        if (ix < 0 || ix >= InX || iy < 0 || iy >= InY) continue;

        for (int bc=0; bc<BS; bc++) // loop over batch
            for(int ic=0; ic<InC; ic++) // loop over inp feature maps
                for(int fc=0; fc<KnC; fc++) // loop over filters
                    OUT[bc][oy][ox][fc] += INP[bc][iy][ix][ic] * K[fy][fx][ic][fc];
    }
}
```


2D convolution backward loop (1):

- As for linear layers, we want to compute two things:
 - Gradient of loss wrt. input
 - Gradient of loss wrt. parameters




- Gradient of Loss wrt. parameters is...
 - a matrix of same shape as parameters
 - obtained by multiplying input and gradient of output:

```
for (int iy = 0; iy < InY; iy++) // loop over input Y
for (int ix = 0; ix < InX; ix++) // loop over input X
{
    for (int fy = 0; fy < KnY; fy++)
    for (int fx = 0; fx < KnX; fx++)
    {
        const int oy = ( iy + Py - fy ) / Sy;
        const int ox = ( ix + Px - fx ) / Sx;
        //padding: skip addition if outside input boundaries
        if (oy < 0 || oy >= OpX || ox < 0 || ox >= OpY) continue;

        for (int bc=0; bc<BS; bc++) // loop over batch
            for (int ic = 0; ic < InC; ic++) // loop over inp feature maps
                for (int fc = 0; fc < KnC; fc++) // loop over filters
                    GK[fy][fx][ic][fc] += ERR[bc][oy][ox][fc] * INP[bc][iy][ix][ic];
    }
}
```


2D convolution backward loop (2):

- As for linear layers, we want to compute two things:
 - Gradient of loss wrt. input
 - Gradient of loss wrt. parameters

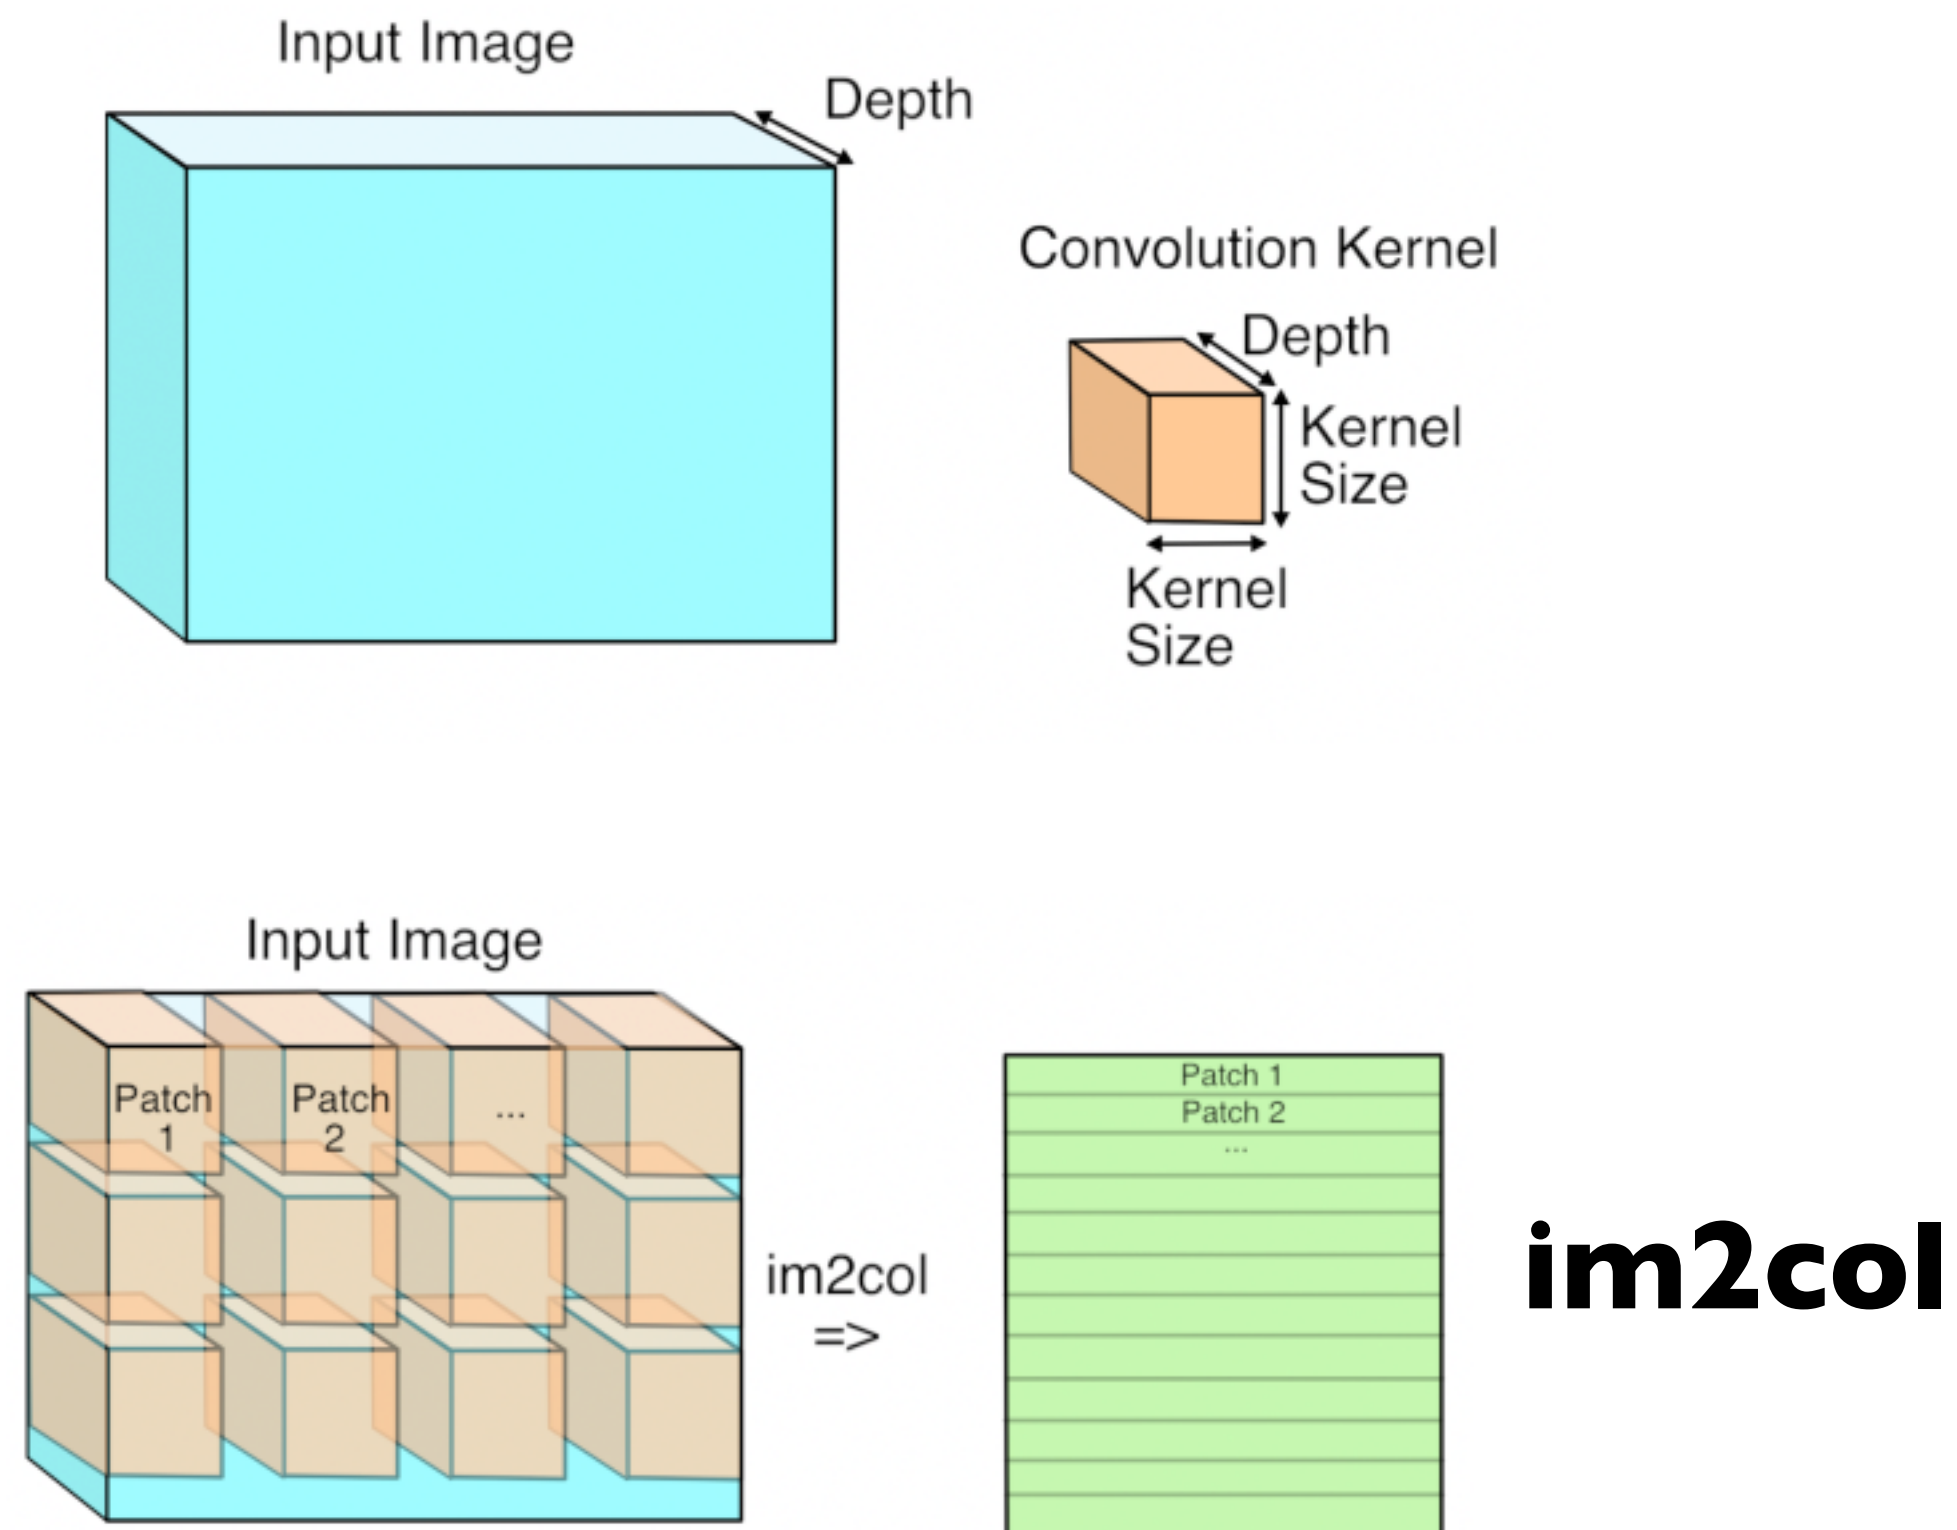


- Gradient of Loss wrt. input is...
 - a matrix of same shape as input
 - obtained by multiplying filter parameters and gradient of output:

```
for (int iy = 0; iy < InY; iy++) // loop over input Y
for (int ix = 0; ix < InX; ix++) // loop over input X
{
    for (int fy = 0; fy < KnY; fy++)
    for (int fx = 0; fx < KnX; fx++)
    {
        const int oy = ( iy + Py - fy ) / Sy;
        const int ox = ( ix + Px - fx ) / Sx;
        //padding: skip addition if outside input boundaries
        if (oy < 0 || oy >= OpX || ox < 0 || ox >= OpY) continue;


        for (int bc=0; bc<BS; bc++) // loop over batch
            for (int ic = 0; ic < InC; ic++) // loop over inp feature maps
                for (int fc = 0; fc < KnC; fc++) // loop over filters
                    GINP[bc][iy][ix][ic] += ERR[bc][oy][ox][fc] * K[fy][fx][ic][fc];
    }
}
```

Exercise: how do we GEMM this?



- ML libraries generally split convolution in two steps: “im2col” and GEMM
- Each convolution is performed by “dotting” a 3D patch of the image with the filter
- im2col copies these patches and organises them into a matrix of size
 $(BS * OpY * OpX) \times (KnY * KnX * InC)$
- One row for each x and y pixel of the batch
- Column of size of the inner prod with one filter
- im2col generally produces an output which occupies more memory than the input img

Exercise: how do we GEMM this?




- Now that memory is aligned, compute output with GEMM:

(mini-batch of output images)


$$[(BS * OpY * OpX) (KnC)] = [(BS * OpY * OpX) (KnY * KnX * InC)] * [(KnY * KnX * InC) (KnC)]$$

(output of im2col) (filter parameters)



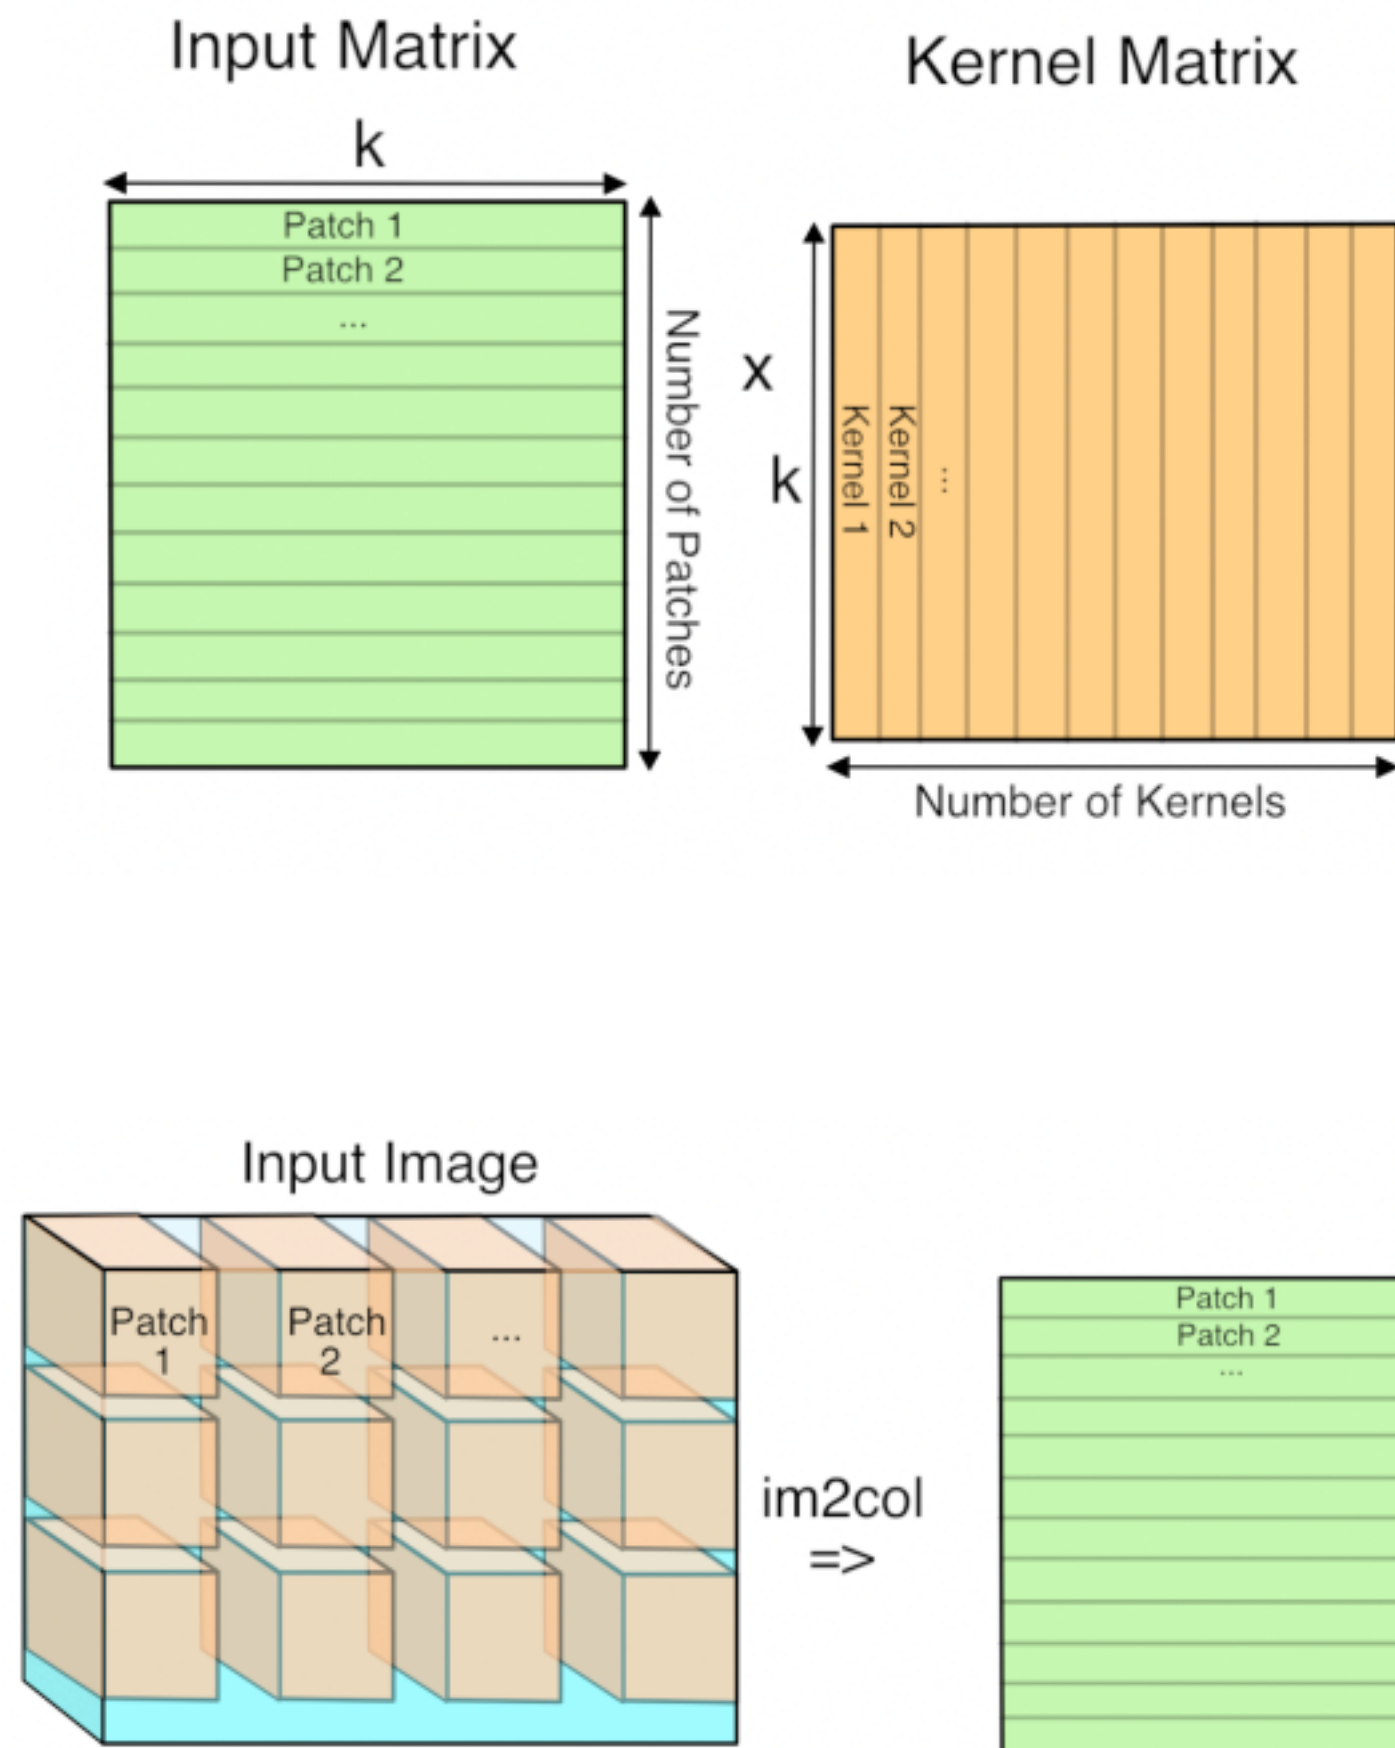
im2col

- Likewise, we split conv layer into two layers:
 - im2col layer
 - conv2d (gemm) layer



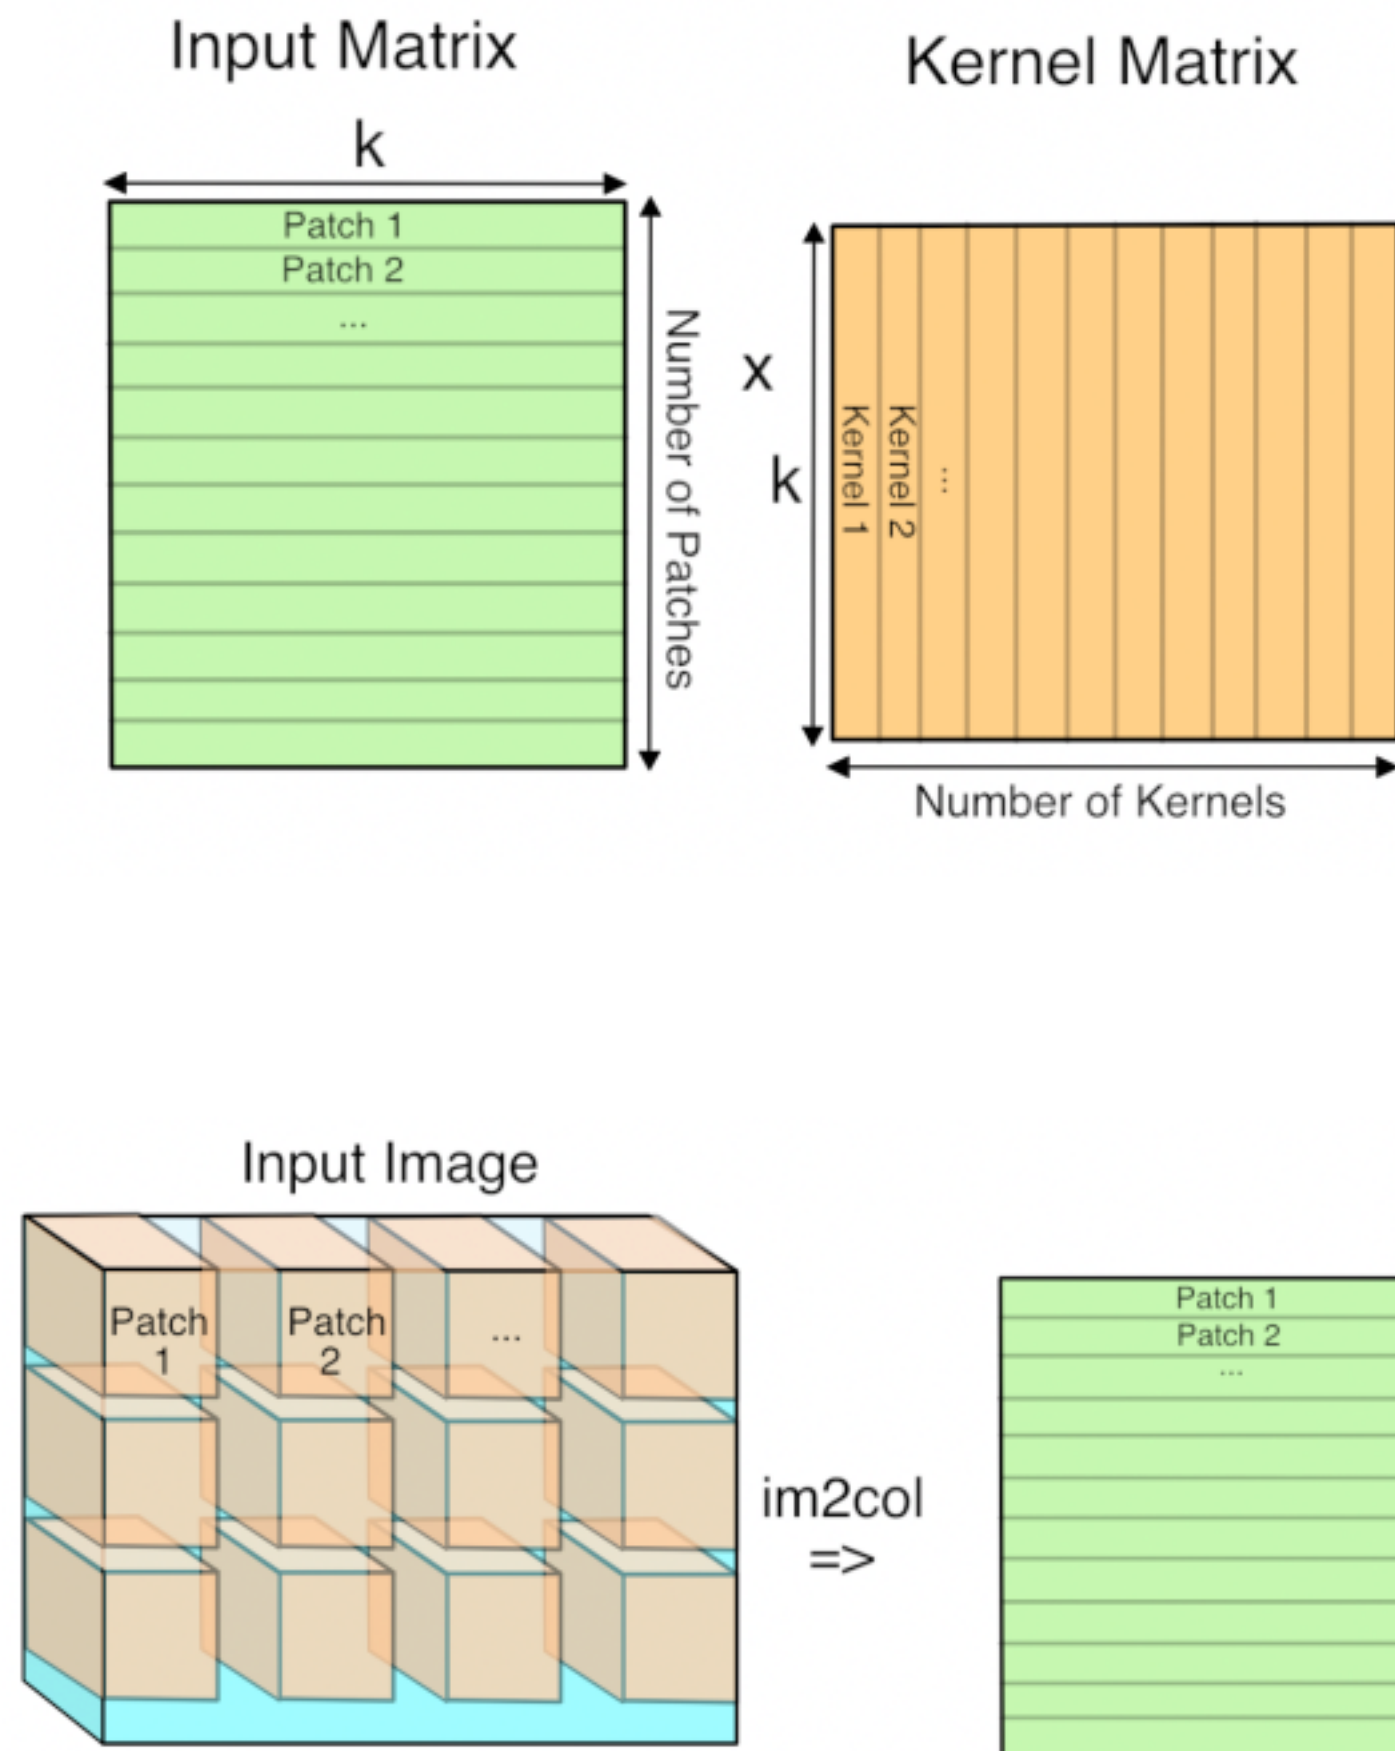
GEMM

Exercise: how do we backprop this?



- Backprop will begin at output layer and proceed to input
- Therefore, for convolution, first we backprop the conv2d layer and then we backprop the im2col layer
- conv2d layer now behaves like linear layer; implement the operation like in exercise 6:
 - dL/dW , dL/dB , $dL/dInp$ are performed by GEMM and loops
- backprop of im2col layer is col2im

Exercise: how do we backprop im2col?



- Backprop of conv2d computes $D = dL / \mathbf{dOutput}$ of Im2Col
- matrix $D : [(BS * OpY * OpX) (KnY * KnX * InC)]$
- As in the im2col step, multiple entries of this matrix are associated with each input pixel
- Backprop of im2col computes $E = dL / \mathbf{dInput}$ of Im2Col
- matrix $E : [(BS) (InY * InX * InC)]$
- Loop through E . For each pixel in (InY, InX) add up the errors contained in D for the corresponding (OpY, OpX, KnY, KnX)
- Mapping between (InY, InX) and the various (OpY, OpX, KnY, KnX) is the same as slide "2D convolution backward loop"

Exercise: a major help

- This is of course not the only way to iterate through arrays
- It sure is convenient
- Only works if array sizes are known at compile time

```
void Im2Mat(const int BS, const Real*const lin_inp, Real*const lin_out) const  
{
```

Input is some linear memory allocation

```
// Convert pointers to a reference to multi dim arrays for easy access:  
// 1) INP is a reference: i'm not creating new data  
// 2) The type of INP is an array of sizes [???][InY][InX][InC]  
// 3) The first dimension is the batchsize and is not known at compile time  
// 4) Because it's the slowest index the compiler does not complain  
// 5) The conversion should be read from right to left: (A) convert lin_inp  
// to pointer to a static multi-array of size [???][InY][InX][InC]  
// (B) Return the reference of the memory space pointed at by a.  
const Real (& INP )[][InY][InX][InC] =  
    * (Real(*)[][InY][InX][InC]) lin_inp;  
// (B) ( A )
```

Convert to C array

(works only if sizes are known at compile time, except for the slowest index which can be unknown)

```
??? = INP[bc][iy][ix][ic];
```

Easy and clear access

Overall objective of the exercise

- Learn to classify MNIST digits

Learn to label each sample:

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2	2
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3
4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4	4
5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5	5
6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6	6
7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7	7
8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8	8
9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9	9

Output of the network is the probability of each image of being a certain digit



CNN



{
0.02
0.03
0.01
0.01
0.70
0.02
0.02
0.01
0.06
0.12
}

Probability that image is

0
1
2
3
4
5
6
7
8
9

Changes from Ex06

- For this exercise we use the Adam optimizer (Kingma & Ba, 2014)
- The non-linearity after each layer is a “Leaky” rectifier linear unit (LReLU):

$$\mathbf{f}(\mathbf{x}) = \mathbf{x} > 0 ? \mathbf{x} : 0.1 \mathbf{x}$$

- Output layer is a SoftMax: $f(x_i) = \frac{\exp x_i}{\sum_{j=1}^{10} \exp x_j}$
- Sum of outputs is one: they behave like probabilities for each digit
- Loss function is the Cross Entropy: $H(\tilde{f}, f) = - \sum_{i=1}^{10} \tilde{f}(x_i) \log f(x_i)$
- Measure of dissimilarity between probability distributions
- Minimized if the distributions are identical
- Target distribution $\tilde{f}(x)$ are the labels. What is probability that an image is each digit?
Well, this is a 4, so:



$$\tilde{f}(x) = \{0, 0, 0, 0, \mathbf{1}, 0, 0, 0, 0, 0\}$$