

Exercise 10 — MPI I/O, Weak Scaling

Question 1: Data Compressor

Concepts:

- MPI I/O: Collective/Non-Collective operations
- File offsets: Determine correct byte offsets to write a file in parallel

Question 2: Weak Scaling of Data Compressor

Concepts:

- Scaling analysis on paper
- Understand scaling behavior of parallel code

Question 1 — MPI I/O for Data Compressor

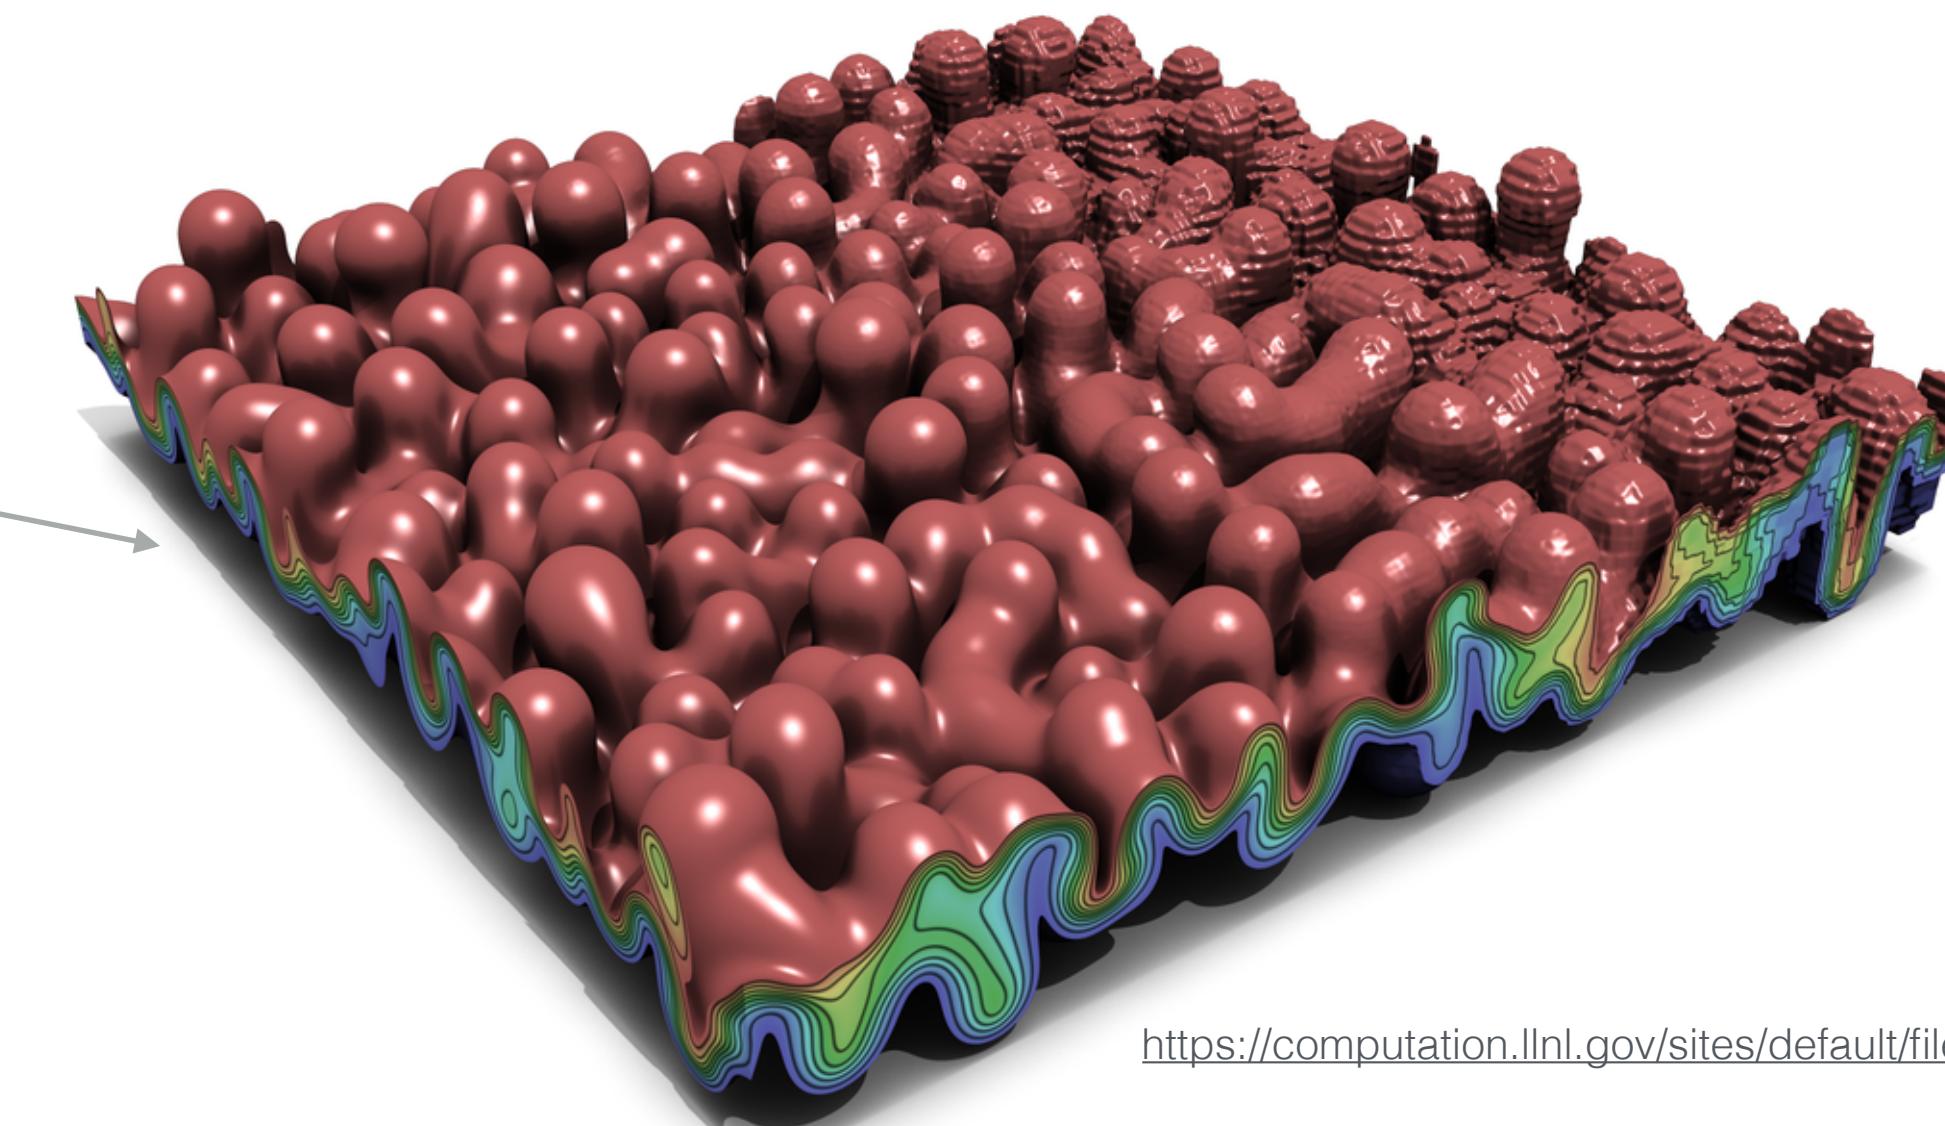
- Code performs floating point data compression for 2D input data
- The compressor is already implemented
- You focus on implementing a parallel file format with the compressed data
- Your code must be able to read and write a file format that you define in parallel using MPI file operations
- We will work with a 2D image file as input to the compressor

Question 1 — MPI I/O for Data Compressor

Data compression:

- Important for HPC computing
- Examples: Reduce disk footprint; Reduce communication overhead
- Special interest in scientific computing: Floating point compression
- Near-lossless or lossy
- We will use the [zfp](#) (lossy) floating point compressor

Near-lossless Compression:
Information is preserved



Lossy Compression: Information
is lost during compression/
decompression cycle

Question 1 — MPI I/O for Data Compressor

Get zfp library and image data before you start:

`make setup`

This will do:

1. Download zfp library and compile it
2. Download the input image `cyclone.bin.gz`

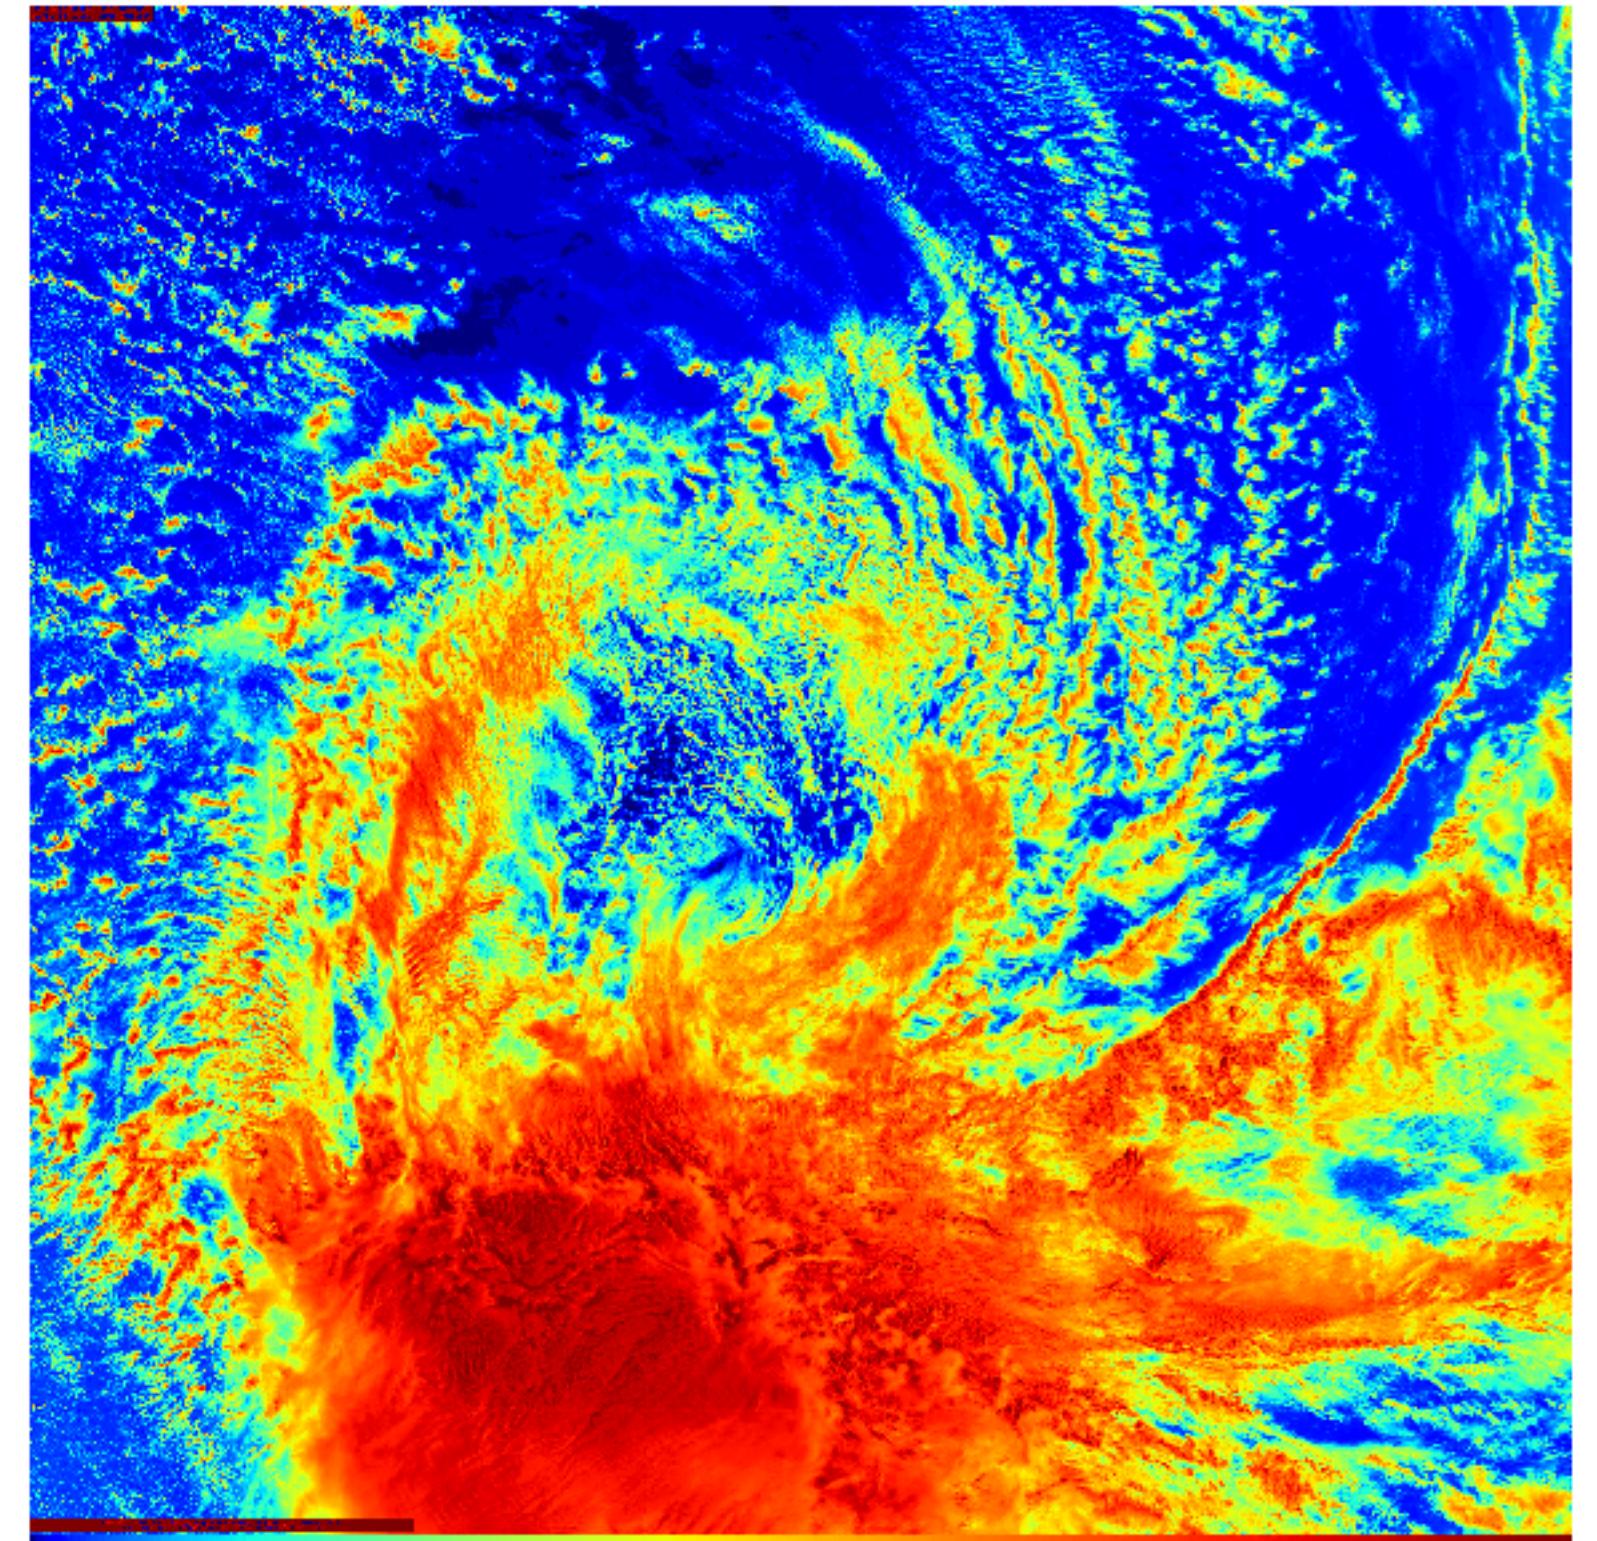
You can compile the code with:

`make`

`mpirun -n <p> ./mpi_float_compression <tol> 4096 cyclone.bin.gz`

Number of MPI processes.
The image dimension 4096
must be evenly divisible by p

Compression tolerance. 0 results in
near-lossless compression (within
machine precision)



Input image dimension (square
image) and image filename

Question 1 — MPI I/O for Data Compressor

Part a.) Implement file protocol for data compressor

You have to complete the function `write_data`:

- The function will write your custom file format for the compressed data (similar to a `.zip` file for example)
- You need to store meta data in the file which you need when you read the file at a later time
- Each MPI rank writes the compressed buffer obtained from `zfp` into the file
- You have to use `MPI_File...` routines. You can use collective or non-collective calls.

Question 1 — MPI I/O for Data Compressor

Part a.) Implement file protocol for data compressor

You can find the meta data that you need to store in the file in the FileHeader and BlockHeader structures:

Required once

```
1 struct FileHeader  
2 {  
3     double tol;           ← Compression tolerance  
4     size_t Nx, Ny, Nb;   ← Dimensions of 2D data  
5 };
```

Nx, Ny, Nb → Number of compressed blocks stored in the file

Required once for each compressed block

```
1 struct BlockHeader  
2 {  
3     size_t start, compressed_bytes, bufsize;  ← Number of bytes in compressed block  
4 };
```

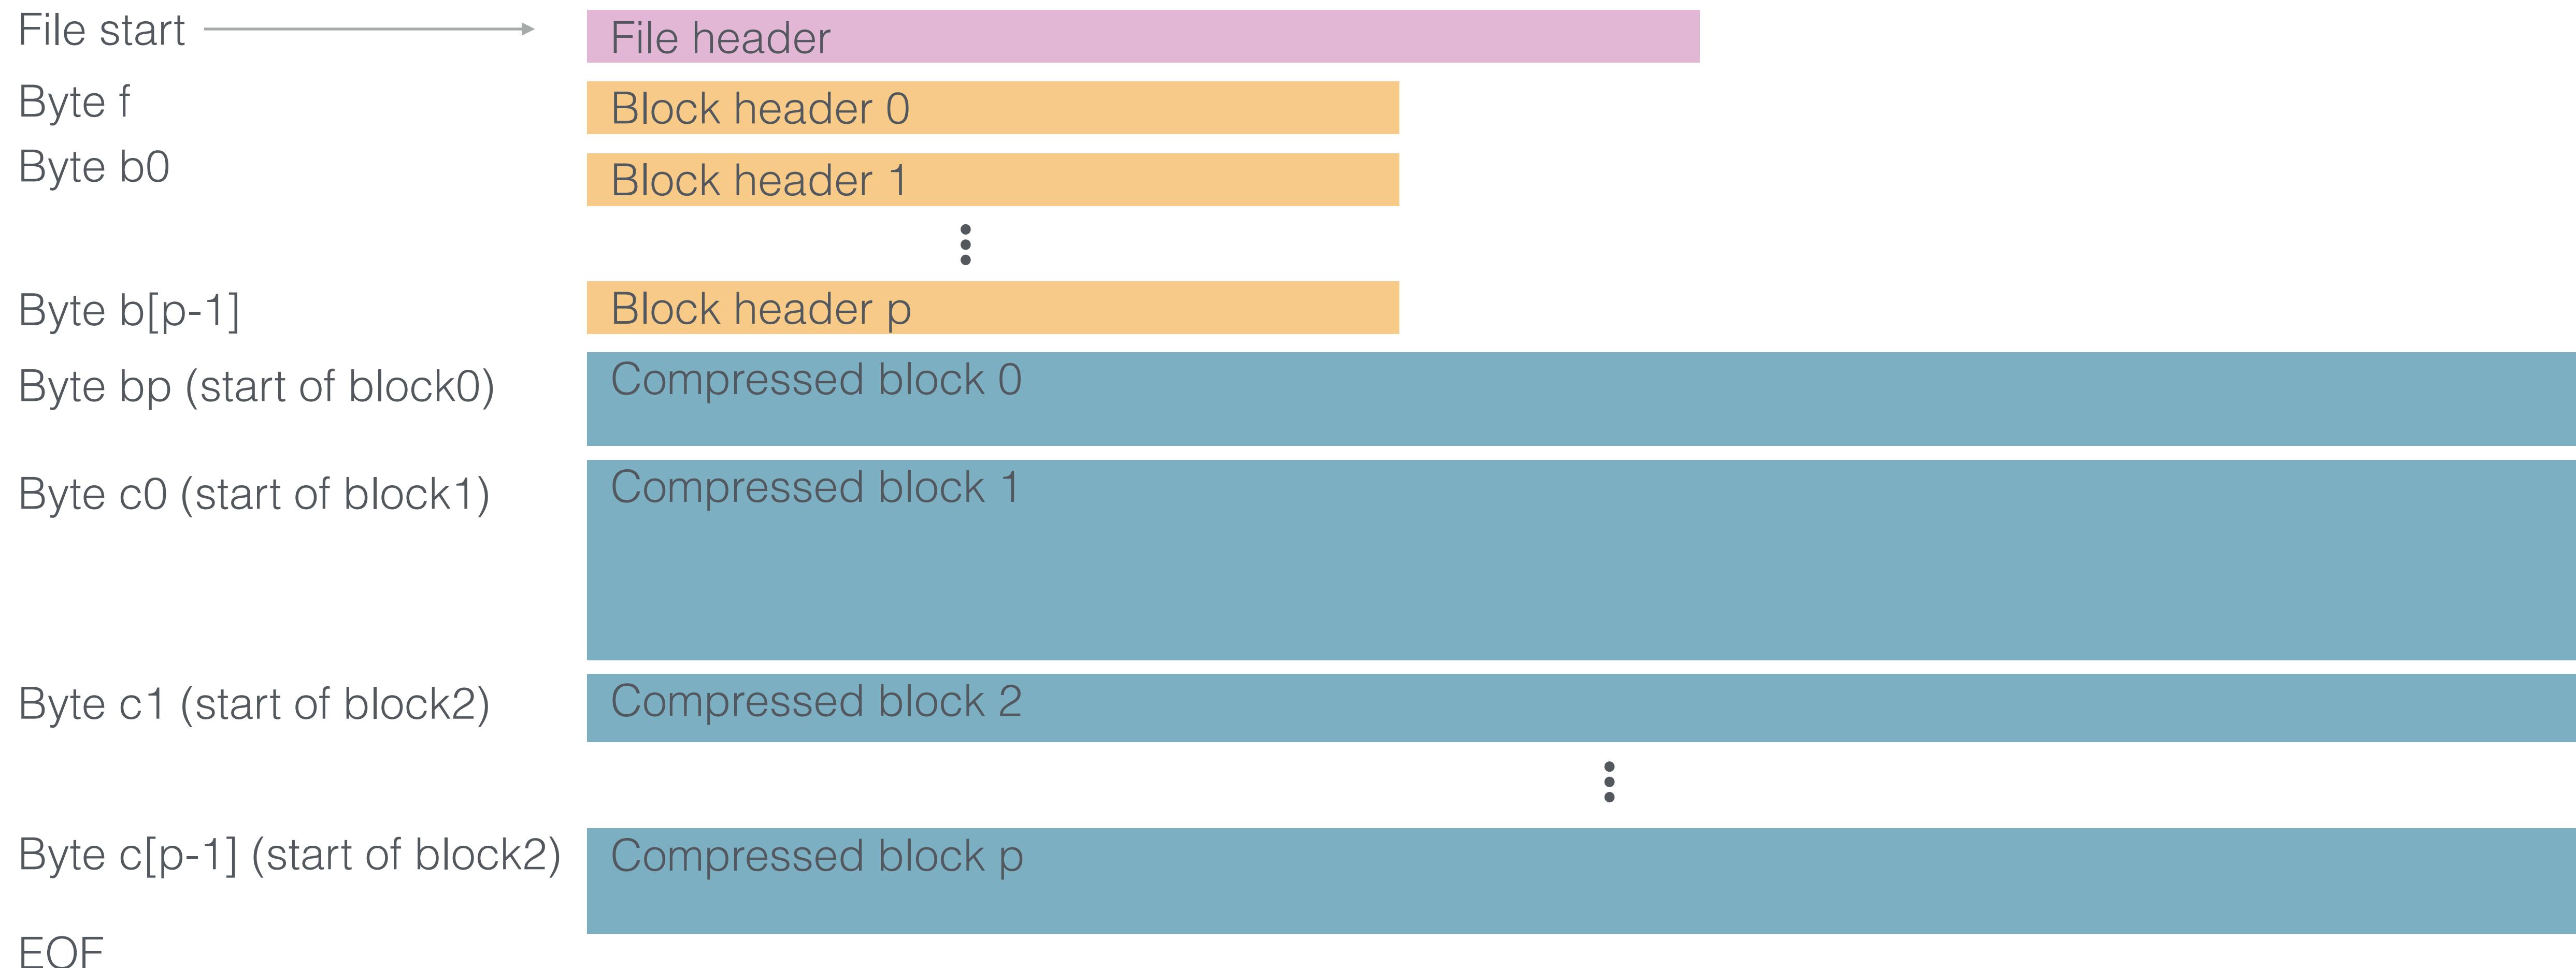
start → Byte offset in the file for compressed block

bufsize → Size of work buffer needed to decompress the data

Question 1 – MPI I/O for Data Compressor

Part a.) Implement file protocol for data compressor

You are free to implement the file protocol yourself. One way to do it:



Question 1 – MPI I/O for Data Compressor

Part b.) Implement the inverse operations of part a.)

You have to complete the function `read_data`:

- The function will read your custom file format in order to decompress the stored blocks
- You need to read the meta data in the file and pass it to the decompressor (pass by reference, see function signature)
- For simplicity, the same number of MPI ranks is used for reading and each rank again reads one compressed block and decompresses that block.
- The decompressed data is written to a `.bin.gz` file and can be compared to the original image with the `print_png.py` Python script.

Question 1 – MPI I/O for Data Compressor

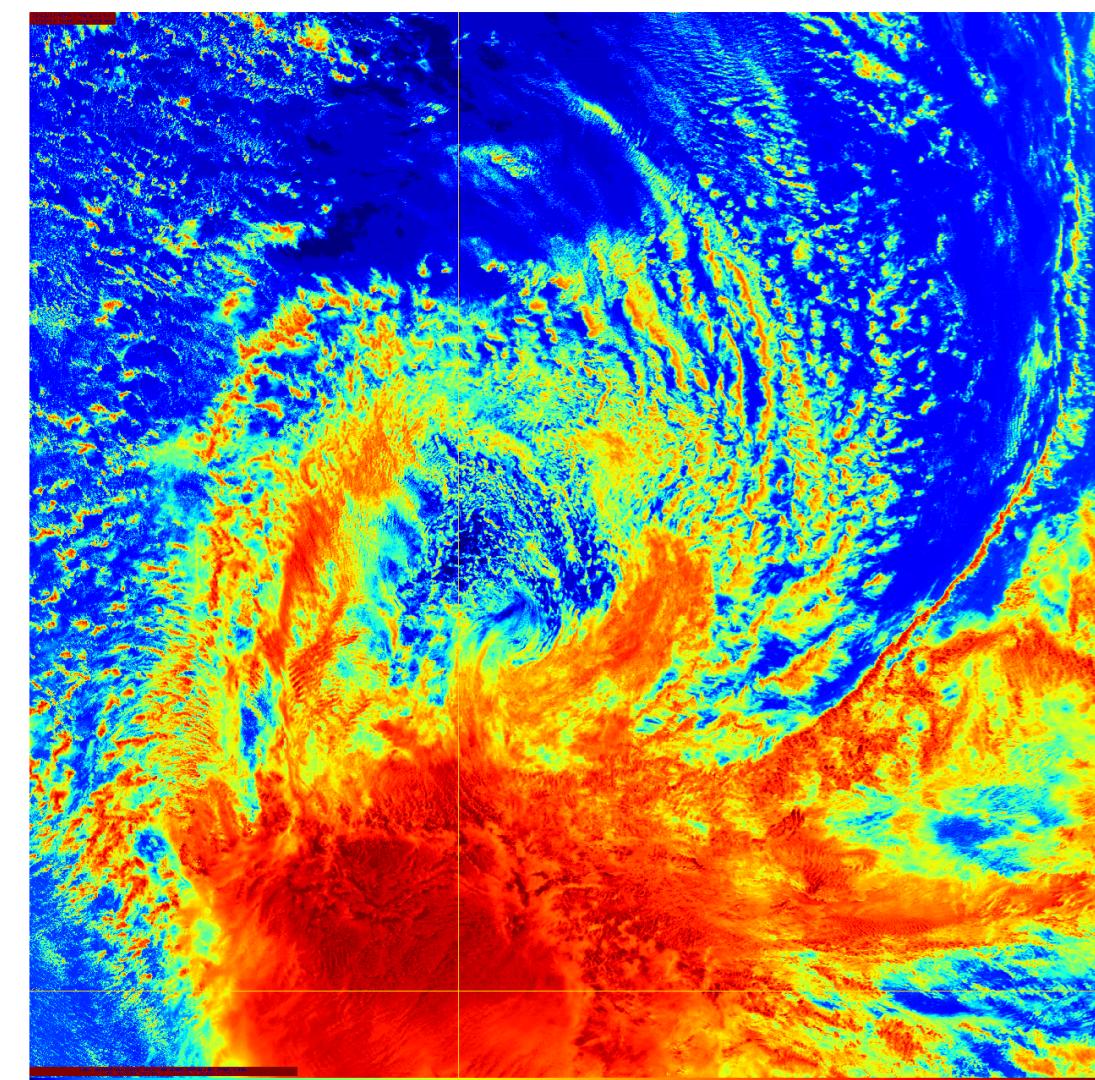
Part b.) Implement the inverse operations of part a.)

To simplify the post-processing with Python, you can use

```
./run.sh <tol>
```

Script does the following:

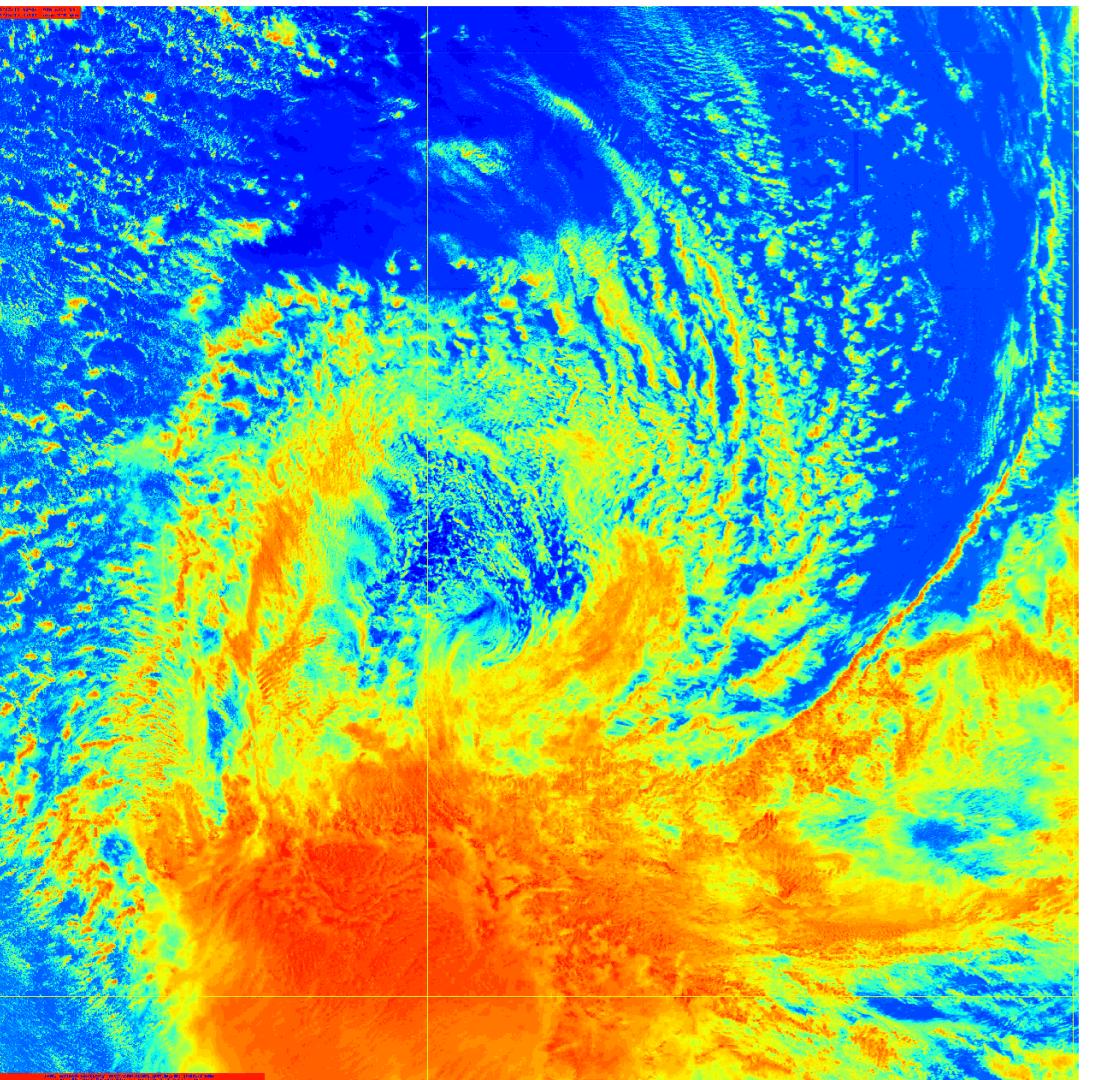
1. Source the build environment
2. Compile the code
3. Run it with MPI
4. Run the Python post-processing



128 MB

Compression rate (zfp): 37.7

tol=200



3.4 MB

Compression rate (zip): 7.5

Question 2 — Weak Scaling

- Pen and Paper
- You are given a table with execution time of the compressor in Question 1
- Measurements correspond to different number of MPI processes and different (square) image sizes N
- Some of the measurement data is not suitable for this problem
- For the weak scaling the problem size *per process* remains constant
- Identify the correct processor/image size pairs and compute the weak efficiency and draw these points in the blank diagram

$$E_w = \frac{t_1}{t_p}$$